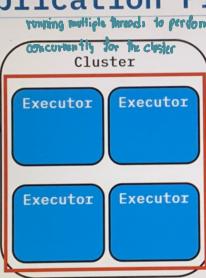


# Apache Spark Architecture

## Spark Application Processes

Can be run on a cluster node or another machine as a client

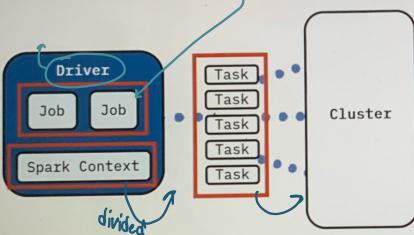


A Spark Application has two main processes:

- Driver Program: a single process that creates work for the cluster and sends it to a cluster work independently
- Executors: multiple processes throughout the cluster that do the work in parallel

## Spark Jobs

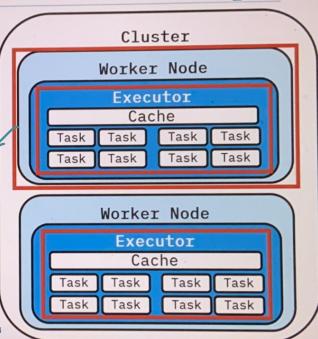
Creates work from the user called "Job"



- Jobs are computations that can be executed in parallel
- The Spark Context divides Jobs into Tasks to be executed on the Cluster

## Spark Executors run tasks from a job

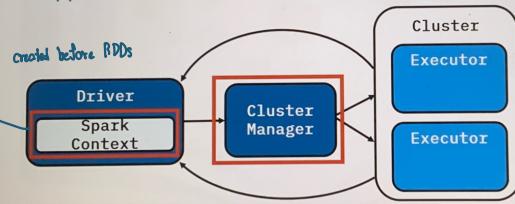
- A Worker is a cluster node that can launch executor processes to run tasks
- Each Executor is allotted a set number of cores that each run one task at a time
- Increasing Executors and cores increases cluster parallelism
- Ideally, limit Executor x core combinations to total cores per node - the example here shows 1 x 8 cores



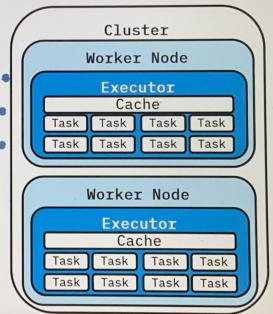
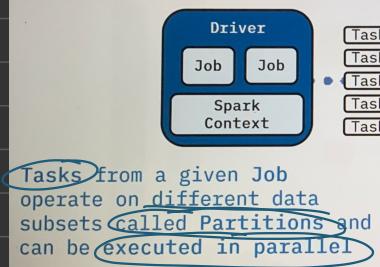
## What is the Spark Context

Spark Context: starts when the application launches and must be created in the driver

- Communicates with the Cluster Manager
- any data frame or RDD created under the context
- Is defined in the Driver, with one Spark Context per Spark Application



## Spark Tasks



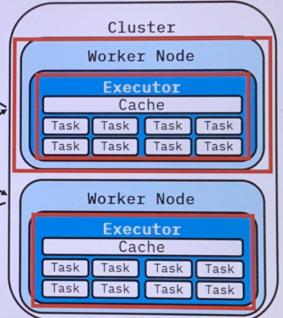
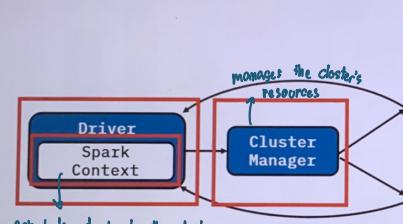
## Why Shuffle Data

A Shuffle is:

- Costly - requiring data serialization, disk and network I/O
- Necessary when an operation requires data outside the current partition of the task
- How Spark re-distributes the dataset across the cluster

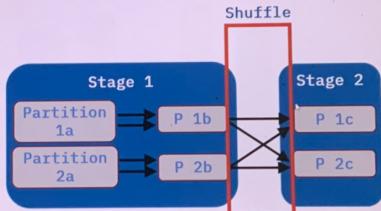
Recap

## Apache Spark Architecture

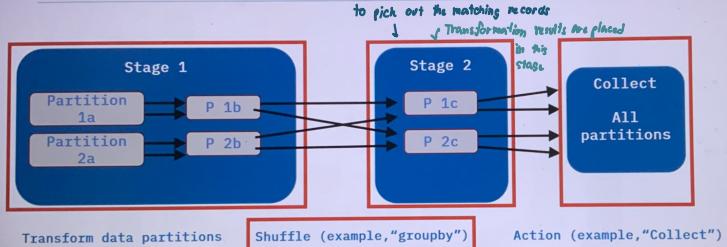


## Spark Stages and Shuffles

- A Stage is a set of tasks within a job that can be completed on the current local data partition
- A Shuffle marks the boundary between Stages
- Stages connect to form a dependency graph



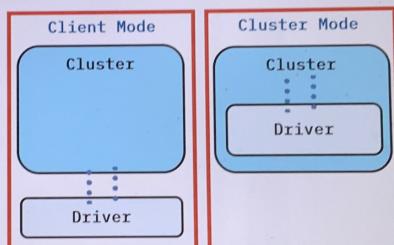
## Shuffle Example



## Driver Deploy Modes

There are two deploy modes:

- Client Mode - the application submitter launches the driver process outside the cluster
- Cluster Mode - the framework launches the driver process inside the cluster



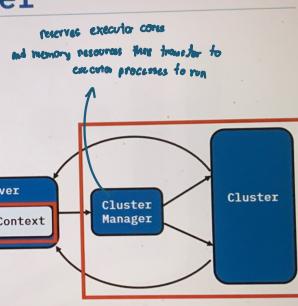
# Overview of Apache Spark Cluster Modes

## Spark Cluster Manager

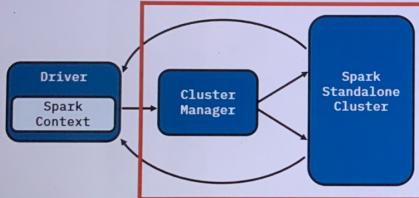
### Spark Cluster Manager:

- Communicates with a particular cluster to acquire resources for the running application
- Runs as a service outside the application and abstracts the cluster type

*Note: An App. is running, the Spark context creates tasks and communicates to the cluster manager what resources are needed*



## Why use Spark Standalone



- Spark Standalone benefits are:
- Built into the Spark installation, so no additional dependencies to deploy
  - Fastest way to setup a Spark cluster and get running

## Types of Cluster Managers

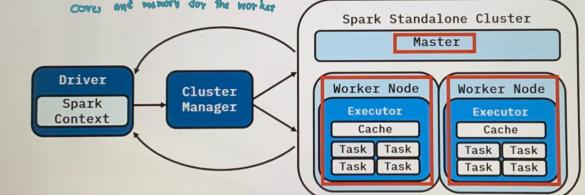
Spark supports the following cluster managers:

- Spark Standalone - comes with Spark, best for setting up a simple cluster
- Apache Hadoop YARN (Yet Another Resource Negotiator) - cluster manager from the Hadoop project
- Apache Mesos - general-purpose cluster manager with additional benefits
- Kubernetes - open-source system for running containerized applications

## Spark Standalone components

Spark Standalone has two main components:

- **Workers** - run an Executor process to receive tasks
  - **Master** - connects and adds workers to the cluster and keeps track of them
- Note: If stuck together, do not reserve all the node's cores and memory for the worker*



## How to set up Spark Standalone

1. Start the Master to output the Master URL

```
$ ./sbin/start-master.sh
```

2. Start Worker(s) with the Master URL

```
$ ./sbin/start-slave.sh spark://<master-spark-URL>:7077
```

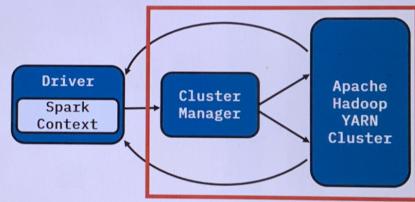
3. Launch Spark application on cluster by specifying the Master URL

```
$ ./bin/spark-submit \
--master spark://<spark-master-URL>:7077 \
<additional configuration>
```

## Why use Apache Hadoop YARN

When choosing Apache Hadoop YARN, consider that it:

- Is general-purpose
- Supports many other big data ecosystem frameworks
- Requires its own configuration and setup
- Has dependencies, making it more complex to deploy than Spark Standalone



## How to run Spark on existing YARN cluster

1. Specify the Master option `--master YARN` with `spark-submit`

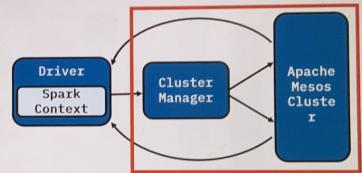
```
$ ./bin/spark-submit \
--master YARN \
<additional configuration>
```

2. Spark will automatically connect with YARN using Hadoop configuration

## Why use Apache Mesos

Apache Mesos Cluster Managers can run Spark with other benefits, such as making partitioning:

- Scalable between many Spark instances
- Dynamic between Spark and other big data frameworks



To configure Spark to run on an Apache Mesos cluster manager see:  
<https://spark.apache.org/docs/latest/running-on-mesos.html>

## Why use Spark on Kubernetes

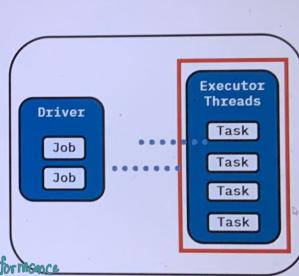
Kubernetes cluster managers can run containerized applications, making it easier to:

- Automate deployment
- Simplify dependency management
- Scale the cluster

## Why use Local Mode

Spark can also be run in local mode which:

- Does not connect to a cluster, making it easy to get started
  - Runs in the same process that calls 'spark-submit' and uses threads for running executor tasks
  - Can be useful for testing or debugging a Spark application
  - Runs a Spark application locally within a single process which can limit performance
- Note: Local mode isn't designed for optimal performance*



## How to set up Local Mode

To run Spark in local mode, use the master option `--master local[#]` where `#` specifies the number of cores to use

```
# Launch Spark in local mode with 8 cores
$ ./bin/spark-submit \
--master local[8] \
<additional configuration>
```

Use an asterisk '\*' to specify using all available cores

```
# Launch Spark in local mode with all available cores
$ ./bin/spark-submit \
--master local[*] \
<additional configuration>
```

# How to Run an Apache Spark Application

## What is 'spark-submit'

'spark-submit' script:

- Spark's Unified Interface for submitting applications
- Found in the 'bin/' directory
- Easily switches from local to cluster mode by changing a single argument
- Works the same way regardless of cluster manager type or application language

```
$ ./bin/spark-submit <config and options> <application files>
```

## Common 'spark-submit' Options

Option/Setting	Form	Mandatory?
Tell Spark what cluster manager to connect with	'--master'	Yes
Specify the program entry point if using Java or Scala application	'--class <full-class-name>'	Yes
Set how the driver is deployed (Client or Cluster). Default is Client mode	'deploy-mode'	No
Set CPU core and memory usage in executors	'--executor-cores' and '--executor-memory'	No
See available options by cluster manager	'./bin/spark-submit -h'	N/A

## 'spark-submit' Examples

1. Launch Scala SparkPi using a jar, with master YARN. Estimate Pi with 1000 samples

```
# Launching Scala SparkPi to a YARN cluster
./bin/spark-submit \
--class org.apache.spark.examples.SparkPi \
--master YARN \
--path/to/examples.jar \
1000
```

2. Launch Python SparkPi to a Spark Standalone cluster. Estimate Pi with 1000 samples

```
# Launching Python SparkPi to a standalone cluster with master at 207.184.161.138
./bin/spark-submit \
--master spark://207.184.161.138:7077 \
examples/src/main/python/pi.py \
1000
```

## Python Application Dependencies

To manage Python application dependencies, ensure:

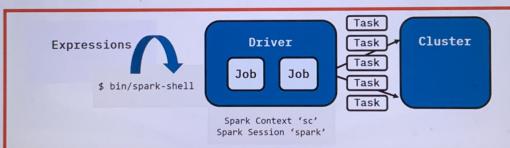
- Cluster nodes can access the required dependencies with same version
- Same Python version is used
- You use the '--py-files' argument so that '.py', '.zip' or '.egg' files can be distributed to the cluster

```
# Launching a PySpark application with dependency on "my_python_package"
$ ./bin/spark-submit <config and options> \
--py-files my_python_package.zip \
my_pyspark_application.py
```

## Spark Shell Environment

For both Scala and Python shells:

- SparkContext is automatically initialized and is available as 'sc'
- SparkSession is automatically available as 'spark' can start working to data immediately
- Expressions are entered into the shell and then evaluated in the driver to become jobs that are scheduled as tasks for the cluster



## Spark Shell Information Provided

- Spark's load log in the log4j-defaults.properties file
- Spark's Web UI address to view information about jobs running in the shell
- Spark Context and Spark Session variables
- Versions of important libraries in use, for example JDK and Scala

```
$ bin/spark-shell
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://192.168.1.33:4040
Spark session available as 'spark'.
Spark session available as 'spark'.
```

version 3.0.1

Using Scala version 2.12.10 (OpenJDK 64-Bit Server VM, Java 1.8.0\_282)
Type in expressions to have them evaluated.
Type :help for more information.
scala>

## Using the 'spark-submit' Script

The 'spark-submit' script will:

- Parse command line arguments/options
- Read additional configuration specified in 'conf/spark-defaults.conf'
- Connect to the cluster manager specified with the '--master' argument or run in local mode
- Transfer application (JARs or Python files) and any additional files specified to be distributed and run in the cluster

## 'spark-submit' Application Files

Final arguments depend on application language:

**Java or Scala**

```
<application-jar-path> <application-args>
Specifies the location of the JAR with your application and dependencies, followed by any arguments specific to the application
```

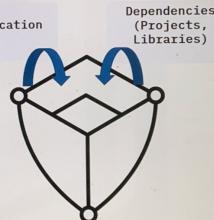
**Python**

```
<application-py-script> <application-args>
Specifies the application python script followed by any arguments specific to the application. Add files with '.py', '.egg' or '.zip' using the '--py-files' argument
```

## Application Dependencies

To manage Spark dependencies:

- Bundle projects or libraries with application so they are accessible to driver and executor processes
- For Java or Scala-based programs, create an uber-JAR with application and dependencies together so it is easy to distribute to the cluster



## Spark Shell

Spark Shell:

- Simple way to learn Spark API
- Powerful tool to analyze data interactively
- Use in local mode or with a cluster, same options as 'spark-submit'
- Can initiate in Scala ('bin/spark-shell') and Python ('bin/pyspark')

**Scala**

```
Run <bin/spark-shell>
```

**Python**

```
Run <bin/pyspark>
```

## Starting Spark Shell - Local Mode

Example: Start Scala Spark-Shell and run local mode by default

```
$ bin/spark-shell
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://192.168.1.33:4040
Spark session available as 'spark' (master = local*, app id = local-1614631042181).
Spark session available as 'spark'.
```

Welcome to

version 3.0.1

Using Scala version 2.12.10 (OpenJDK 64-Bit Server VM, Java 1.8.0\_282)
Type in expressions to have them evaluated.
Type :help for more information.

scala>

## Spark Shell Example - Run Code

```
scala> val df = spark.range(10)
df: org.apache.spark.sql.Dataset[Long] = [id: bigint]
```

```
scala> df.withColumn("mod", expr("id % 2")).show(4)
```

id	mod
0	0
1	1
2	0
3	1

- Launch Scala Spark Shell
- Create distributed DataFrame with column 'id' and 10 values (0-9)
- Add a column that evaluates an SQL expression for modulo of 2 and show first four (4) results