

CRUD Operations

What you will learn



Use Mongo shell to connect a database



Perform basic CRUD operations

Mongo shell

A command line tool that helps you interact with MongoDB

```
> mongo "mongodb://USER:PASSWORD@uri/test"
```

```
Enterprise my-replicaset [primary] test> show dbs
admin          0.000GB
campusManagementDB 0.000GB
local          1.218GB
```

Mongo shell

Set database context

```
> use campusManagementDB
switched to db campusManagementDB
> show collections
staff
students
```

Create

```
> db.students.insertOne({
    "firstName": "John",
    "lastName": "Doe",
    "email": "john.doe@email.com",
    "studentId": 20217484
})
{
    "acknowledged" : true,
    "insertedId" : ObjectId("60424dad68ad40a3490e0db4")
}
```

Create Many

```
> students_list = [
  {"firstName": "Sarah", "lastName": "Jane", "email": "sarah@mail.com", "studentId": 20215124 },
  {"firstName": "Peter", "lastName": "Parker", "email": "peter@web.com", "studentId": 20218873 }]
> db.students.insertMany(students_list)
{
  "acknowledged" : true,
  "insertedIds" : [
    ObjectId("60424ee468ad40a3490e0db5"),
    ObjectId("60424ee468ad40a3490e0db6")
  ]
}
```

Read

```
> db.students.findOne()
{
  "_id" : ObjectId("60421b5c26fe73f38fdd83db"),
  "firstName" : "John",
  "lastName" : "Doe",
  "email" : "johnd@campus.edu",
  "studentId" : 20217484
}
```

Read

```
> db.students.findOne({ "email": "sarah@mail.com" })  
{  
  "_id" : ObjectId("604228bee45df18c77dc2dd7"),  
  "firstName" : "Sarah",  
  "lastName" : "Jane",  
  "email" : "sarah@mail.com",  
  "studentId" : 20215124  
}
```

Read

```
> students.find({ "lastName": "Doe" })  
[ { "_id" : ObjectId("60421b5c26fe73f38fdd83db"), "firstName" : "John", "lastName" : "Doe",  
  "email" : "john.doe@email.com", "studentId" : 20217484 },  
  { "_id" : ObjectId("60421b5c26fe73f38fdd84ab"), "firstName" : "Mark", "lastName" : "Doe",  
  "email" : "mark@mail.com", "studentId" : 20215001 } ]
```

Read

```
> students.count({ "lastName": "Doe" })  
2
```

Replace

```
> student = db.students.findOne({ "lastName": "Doe" })  
> db.student["onlineOnly"] = true  
> db.student["email"] = "johnd@campus.edu"  
> db.students.replaceOne({ "lastName": "Doe" }, student)  
< { "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
```

Update

```
> changes = { "$set" : { "onlineOnly" : true,  
  "email" : "johnd@campus.edu" } }  
  
> db.students.updateOne({ "lastName": "Doe" }, changes)  
< { "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }  
  
> db.students.updateMany({}, { "$set": { "onlineOnly": true } })  
< { "acknowledged" : true, "matchedCount" : 4, "modifiedCount" : 3 }
```

Delete

```
> db.students.deleteOne({ "studentId": 20218873 })  
< { "acknowledged" : true, "deletedCount" : 1 }  
> db.students.deleteMany({ "graduatedYear": 2019 })
```

Indexes

What you will learn



Explain why we need indexes



Perform efficient searches and sorts



Describe how MongoDB stores indexes

Why we need indexes

Indexes help quickly locate data without looking for it everywhere.

- Example: In the British Library, find *Designing Data-Intensive Applications* by Martin Kleppmann
 - Without indexes:** You will need to look through all of the books
 - With indexes:** Go to Computers / Databases / M (for Martin)



When to index

- Most frequent queries

collection name
↓
key value
db.courseEnrollment.find({ "courselid": 1547 })
means store the index in ASC order
db.courseEnrollment.createIndex({ "courselid": 1 })

Indexes in MongoDB

- MongoDB indexes are special data structures
- They store the fields you are indexing
- They also store the location of the document

Everyday index samples

- Telephone book
- Dictionary
- Index at the end of a book

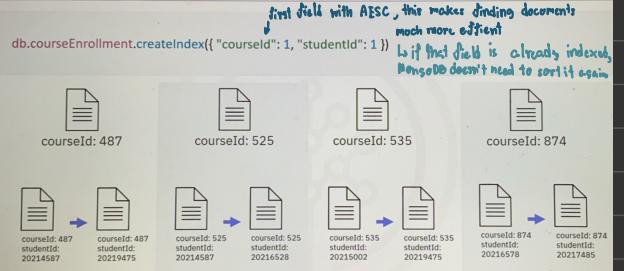


Use indexes to sort

Indexes can help with sorting

change to be a compound index → indexing more than one field.
Once we find all documents that matched 'courselid':
{ "courselid": 1547 } , sort({ "studentId": 1 })
db.courseEnrollment.createIndex({ "courselid": 1, "studentId": 1 })

How MongoDB stores indexes



Aggregation Framework

What you will learn



Explain what an aggregation framework is in MongoDB



Describe how the aggregation framework is built and list its most common stages



Describe use cases for aggregation

Aggregation framework → sometimes referred to as aggregation pipeline

What is the average student score in each course for 2020?

```
> db.courseResults.aggregate([
  { $match: { "year": 2020 } },
  { $group: { "id": "$courseId", "avgScore": {
    $avg: "$score" } } }
])
```

After first avg score

Course results

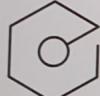
```
courseResults = [
  { _id: 1, courseId: 1, year: 2019, score: 95 },
  { _id: 2, courseId: 1, year: 2019, score: 84 },
  { _id: 3, courseId: 2, year: 2019, score: 91 },
  { _id: 4, courseId: 1, year: 2020, score: 95 },
  { _id: 5, courseId: 1, year: 2020, score: 100 },
  { _id: 6, courseId: 2, year: 2020, score: 96 },
  { _id: 7, courseId: 2, year: 2020, score: 89 },
  { _id: 8, courseId: 2, year: 2020, score: 73 },
  { _id: 9, courseId: 2, year: 2021, score: 92 },
  { _id: 10, courseId: 1, year: 2021, score: 77 },]
```

Aggregation stages

```
> db.courseResults.aggregate([
  { $match: { "year": 2020 } },
  { $group: { "_id": "$courseId", "avgScore": { $avg: "$score" } } }
])
```

→ separated documents into groups according to a group key
→ the id field is not grouped
→ hold avg value

Common aggregation stages



- \$project → reshapes documents by including/excluding fields
- \$sort
- \$count count in that stage
- \$merge move output to collection

Applying \$project

```
students = [
  { "_id": 1, "firstName": "Sophie", "lastName": "Foucher", "city": "Quebec" },
  { "_id": 2, "firstName": "Stafford", "lastName": "von Hagt", "city": "Terezin" },
  { "_id": 3, "firstName": "Olivero", "lastName": "Liggett", "city": "Guaranda" }]
```

Applying the \$project stage

```
db.students.aggregate([
  {
    $project: { firstName: 1, lastName: 1 }
  }
])
```

Applying \$project: Resulting outcomes

```
[{"_id": 1, "firstName": "Sophie", "lastName": "Foucher"}, {"_id": 2, "firstName": "Stafford", "lastName": "von Hagt"}, {"_id": 3, "firstName": "Olivero", "lastName": "Liggett"}]
```

→ items included unless explicitly excluded

Applying the \$sort stage: Sorting order options

\$sort

```
db.students.aggregate([
  {
    $project: { firstName: 1, lastName: 1 }
  },
  {
    $sort: { lastName: -1 } → sorting by lastName
  }
])
```

→ sorting by lastName in DESC

Applying \$sort stage: The resulting output

```
[{"_id": 2, "firstName": "Stafford", "lastName": "von Hagt"}, {"_id": 3, "firstName": "Olivero", "lastName": "Liggett"}, {"_id": 1, "firstName": "Sophie", "lastName": "Foucher"}]
```

Applying the \$count stage

\$count

```
> db.students.aggregate([{$count: "totalStudents"}])
{ "totalStudents": 3 }
```

Applying \$match

```
courseResults = [
  { _id: 1, courseId: 1, year: 2019, score: 95 }, → $merge
  { _id: 2, courseId: 1, year: 2019, score: 84 },
  { _id: 3, courseId: 2, year: 2019, score: 91 },
  { _id: 4, courseId: 1, year: 2020, score: 95 },
  { _id: 5, courseId: 1, year: 2020, score: 100 },
  { _id: 6, courseId: 2, year: 2020, score: 96 },
  { _id: 7, courseId: 2, year: 2020, score: 89 },
  { _id: 8, courseId: 2, year: 2020, score: 73 },
  { _id: 9, courseId: 2, year: 2021, score: 92 },
  { _id: 10, courseId: 1, year: 2021, score: 77 }]
```

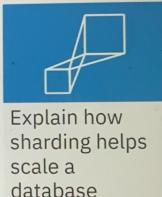
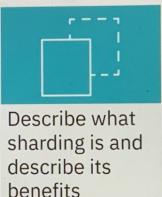
Aggregation use cases

- Reporting: Track student progress
- Analysis: Determine the average product sales by country

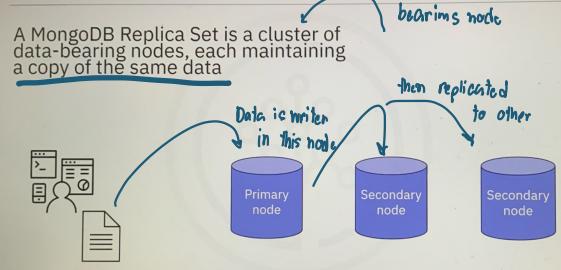


Replication and Sharding

What you will learn

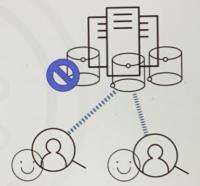


Replication

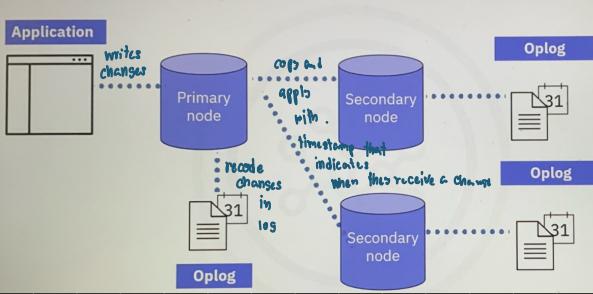


Benefits of replication

- Redundancy → if one node failed, it still runs from other copied data
- High availability
- Fault tolerance



The replication process



How an Oplog works

The command, db.students.deleteMany({}) displays as an individual operation in the Oplog

```
{ "op" : "d", "ns" : "campusManagement.student", "o" : { "_id" : 1 } }
{ "op" : "d", "ns" : "campusManagement.student", "o" : { "_id" : 2 } }
{ "op" : "d", "ns" : "campusManagement.student", "o" : { "_id" : 3 } }
```

How an Oplog works

Example: A database of student information:

```
{"_id":1,"firstName":"Sophie","lastName":"Foucher","city":"Quebo"} ...
{"_id":2,"firstName":"Stafford","lastName":"von Hagt","city":"Terezin"} ...
{"_id":3,"firstName":"Olivero","lastName":"Liggett","city":"Guaranda"}
```

You can delete the preceding data using the statement:
db.students.deleteMany({})

What is Sharding?

When data is growing beyond hardware capacity:

- Scale vertically (bigger, faster hardware) → sometimes it's not feasible
- Scale horizontally (partition the data)



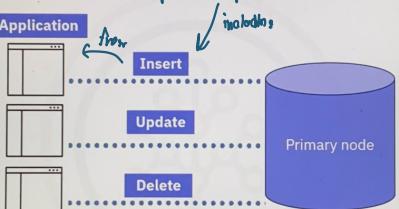
Benefits of Sharding

- Helps with increased throughput and capacity
- Can help adhere to legal requirements for data storage



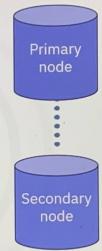
The Primary node and elections

Only node that can accept write operations



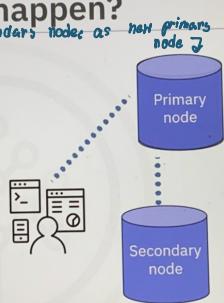
Requirements to become Primary

- MongoDB's voting system ensures only one member becomes a primary node
- Member nodes that are eligible have minimal replication lag and have the most recent data
- The member node receiving the majority of the votes becomes the new Primary.



Why does an election happen?

- The current primary node becomes unavailable
- A new replica set is initialized and needs to choose its initial primary
- A database administrator initiates a manual failover for maintenance or upgrades



Accessing MongoDB from Python

What you will learn



Explain what the MongoClient is



Perform basic CRUD operations using Python

MongoClient

A class that helps you interact with MongoDB

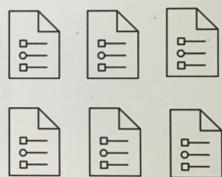
```
from pymongo import MongoClient
uri = "mongodb://USER:PASSWORD@uri/test"
client = MongoClient(uri) ← should do once in our code
campusDB = client.campusManagementDB ← object points
students = campusDB.students ← point collection
```

Create Many

```
students_list = [
    {"firstName": "Sarah", "lastName": "Jane",
     "email": "sarah@mail.com", "studentId": 20215124},
    {"firstName": "Peter", "lastName": "Parker",
     "email": "peter@web.com", "studentId": 20218873}]
students.insert_many(students_list)
```

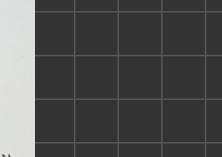
Print read documents

```
from bson.json_util import dumps
cursor = students.find({"lastName": "Doe"})
print(dumps(cursor, indent=4))
```



Update for small changes

```
changes = {"$set": {"onlineOnly": True,
                     "email": "johnd@campus.edu"}}
students.update_one({"lastName": "Doe"}, changes)
students.update_many({}, {"$set": {"onlineOnly": True}})
```



Create

```
students.insert_one({"firstName": "John",
                     "lastName": "Doe",
                     "email": "john.doe@email.com",
                     "studentId": 20217484})
```

Read

```
students.find_one() → return first doc.
students.find_one({"email": "sarah@mail.com"}) → specify
return all doc that equal to ↓
students.find({"lastName": "Doe"}) → read the first one
count ↴
students.count_documents({"lastName": "Doe"})
```

Replace for large changes

```
student = students.find_one({"lastName": "Doe"})
student["onlineOnly"] = True ← create new field
to indicate where student are
student["email"] = "johnd@campus.edu" ← update email
students.replace_one({"lastName": "Doe"}, student)
```

filter ↑ update ↓

Delete

```
students.delete_one({"studentId": 20218873})
students.delete_many({"graduatedYear": 2019})
```